

Proof-of-Continuity: A Temporal Model for Authority Propagation in Distributed Systems and AI Agents

Nicola Gallo

8 July 2026

Abstract

Proof-of-Possession authorization models derive authority from the possession of artifacts such as tokens, credentials, or capabilities. This paper argues that possession is insufficient for discrete execution chains, whether they span multiple services or occur as separated steps within the same machine, because it does not preserve the causal relationship between the origin of a request and the authority exercised at later steps. We introduce Proof-of-Continuity as a minimal authority propagation model in which each execution step must be causally linked to the previous step and may only propagate a non-expansive subset of the authority received from the origin. Under this model, the confused deputy condition cannot be satisfied as valid model behavior: any privilege exercised at a later step must already be present in the origin authority context. This is directly relevant to distributed systems and AI agents, where executors invoke tools and downstream services while holding multiple authority sources, so that the same authority/causality mismatch recurs across service boundaries.

1 Model

We formalize the PIC (Provenance Identity Continuity) Model as follows.

Let P be a finite set of principals, O a set of operations, and R a set of resources. Define a privilege as $(o, r) \in O \times R$.

Each principal p has an associated privilege set:

$$Priv(p) \subseteq O \times R.$$

A principal is used here as an abstract permissioned entity: it may be a human identity established through an identity provider, wallet,

OIDC/OIDC4VC-style flow [13, 14], or decentralized identifier [15]; a workload or machine identity such as those used in WIMSE/SPIFFE-style environments [11, 12]; a role; a service account; or another execution context. What matters for this model is that permissions are associated with it and that it can express an intent within those permissions. When a permissioned entity expresses such an intent, the system creates the origin authority context C_0 for a new execution lineage.

For example, if Bob has the privilege $(\text{read}, \text{Book})$, then two different executions may both involve Bob and the same privilege. They are not therefore the same authority event. One execution may be a valid continuation of an intent Bob expressed within a lineage, while another may be a replay, confused-deputy action, or unrelated execution using the same apparent privilege. The identity and privilege coordinates may be identical; the lineage coordinate distinguishes the authority event.

A request is a message *req* sent between principals and may contain a payload.

Execution is modeled as a causal chain of steps:

$$\pi = \langle (s_0, C_0), (s_1, C_1), \dots, (s_n, C_n) \rangle$$

where s_0 is the origin step and $C_i \subseteq O \times R$ is the authority context carried at step s_i . We use s_i for the formal execution step at hop i ; operationally, a hop is the execution of one such step by an executor that receives, processes, and may propagate a request. Each hop may be performed by a service, workload, function, tool, or agent; the model requires continuity of authority, not propagation of identity.

Definition 1 (PIC Model). *A system enforces Provenance Identity Continuity if, for every execution chain π , each transition preserves causal continuity and monotonic authority restriction:*

$$C_{i+1} \subseteq C_i$$

for all i , and every privilege exercised at hop n is causally linked to the origin via a verifiable chain.

Theorem 1 (PIC Safety). *If the PIC Model holds, no execution step can exercise a privilege not present at the origin.*

Proof. By construction, $C_{i+1} \subseteq C_i$ for all i , so $C_n \subseteq C_0$. Thus, any $(o, r) \in C_n$ must also be in C_0 . Therefore, no privilege can be gained beyond the origin authority context.

2 Threat Model and Scope

We consider execution chains in which an origin request causes one or more later execution steps. Executors may be buggy, confused, or adversarially influenced through their inputs. An executor may hold multiple authority sources — ambient privileges, delegated tokens, bearer credentials, or capabilities — and may attempt to use one unrelated to the request being processed.

The adversary may influence requests, payloads, and parameters passed across the chain, and may cause an executor to select the wrong authority source. The adversary cannot forge the relationship evidence required by *PoR*, nor make a valid continuity checker accept a transition that violates monotonicity. As capability and token models assume their credentials unforgeable, we assume the relationship evidence and the authority-transition mechanism unforgeable; compromise of the validator, the cryptographic root of trust, or the initial authority bootstrap is out of scope.

3 Confused Deputy

Definition 2 (Confused Deputy). *A confused deputy occurs when there exist principals U (user) and D (deputy) and a privilege $(o, r) \in O \times R$ such that:*

1. $(o, r) \notin \text{Priv}(U)$,
2. $(o, r) \in \text{Priv}(D)$,
3. U sends a request req to D ,
4. as a consequence of req , D executes (o, r) .

This definition does not depend on implementation details, only on the mismatch of authority and causality.

Classical example. A FORTRAN compiler `FORT`, installed in the privileged directory `SYSX`, holds ambient authority to write statistics to `(SYSX)STAT`. A user invokes the compiler and provides `(SYSX)BILL` as the debugging output filename. The compiler opens the target for output using its own authority over `SYSX`, thereby overwriting `(SYSX)BILL`. The user never possessed this privilege, but the deputy exercised it as a consequence of the user’s request [1].

4 Proof-of-Possession (PoP)

A token t represents a set of privileges:

$$Priv(t) \subseteq O \times R.$$

PoP Semantics. Possession implies usability: if a principal holds t , it may exercise all $(o, r) \in Priv(t)$.

PoP systems constrain authority by artifact possession, but do not, by themselves, require that the exercised authority be a causal continuation of the authority context of the request.

4.1 Authority Mixing

Assume:

$$(write, r) \notin Priv(U) \tag{H1}$$

$$(write, r) \in Priv(D) \tag{H2}$$

Here U is the user and D is a deputy.

In a possession-based model, a deputy may hold multiple authority sources: its own ambient privileges, delegated tokens, bearer credentials, or other artifacts. The model authorizes an operation by checking whether some possessed authority source grants the requested privilege.

Theorem 2 (PoP admits authority mixing). *In a possession-based model, if an executor may select authority from any artifact or ambient privilege it holds independently of the request authority context, then the confused deputy condition is admissible in the model.*

Proof. Let a request originate from an authority context that does not contain (o, r) . Let the executor hold another authority context, artifact, or ambient privilege that does contain (o, r) .

In a PoP model, authorization depends on the possession of an artifact or privilege granting (o, r) , not on whether that authority is a continuation of the request authority context. Therefore the model permits an execution in which the request causally triggers (o, r) while (o, r) is absent from the request authority context. This is exactly the confused deputy condition.

This applies to any model in which authorization is derived from possession independently of the causal lineage, including JWTs, bearer credentials, and capabilities treated as held artifacts when they may be applied outside

the execution that requested them. The gap is not possession itself but the absence of a temporal binding between the authority used and the execution that caused its use.

5 Proof-of-Continuity (PoC / PIC)

Execution is a causal chain:

$$s_0 \rightarrow s_1 \rightarrow \cdots \rightarrow s_n.$$

Let $C_i \subseteq O \times R$ denote the authority context carried by execution step s_i .

Definition 3 (Proof of Relationship). *PoR(s_i, s_{i+1}) holds iff s_{i+1} is a valid causal continuation of s_i within the same execution lineage. PoR is a relational, single-hop property: it binds one step to its immediate predecessor.*

Definition 4 (Valid transition). *A transition from s_i to s_{i+1} is valid iff it is both causally related and non-expansive:*

$$Valid(s_i, C_i, s_{i+1}, C_{i+1}) \iff PoR(s_i, s_{i+1}) \wedge C_{i+1} \subseteq C_i.$$

The first condition proves causal relationship. The second condition proves monotonic authority restriction. Thus each hop transfers only:

$$C_{i+1} \subseteq C_i. \tag{C1}$$

Definition 5 (Proof of Continuity). *For a chain $\pi = \langle s_0, \dots, s_n \rangle$, PoC(π) holds iff every adjacent transition is valid:*

$$PoC(\pi) \iff \bigwedge_{i=0}^{n-1} (PoR(s_i, s_{i+1}) \wedge C_{i+1} \subseteq C_i).$$

Thus continuity is the transitive composition of relationships: *PoR* proves a single link, *PoC* proves the whole lineage. Relationship is local; continuity is global.

Monotone propagation along a causal chain. Two things propagate together. First, authority attenuates monotonically: write $C_{i+1} \preceq C_i$ for authority *attenuation* — every authority at a hop is a restriction of one held at the previous hop, never an expansion. When the context is a set of atomic privileges this specializes to $C_{i+1} \subseteq C_i$, the form used in our examples.

Second, the hops are bound into a single causal lineage by *PoR*: continuity is a property not of a hop in isolation but of the *chain*. Formally the lineage is a logical causal chain; in an implementation it need not be carried whole — it may be represented by cryptographic commitments, references, hashes, signatures, or summarized proofs that allow each hop to prove its place in the lineage without transporting the entire history.

Heterogeneous operation spaces. The notation assumes a common privilege universe $O \times R$. In heterogeneous systems different hops use different vocabularies; let $O_i \times R_i$ be the universe at hop i . A transition is then governed by a policy translation

$$\mathcal{T}_{i \rightarrow i+1} : \mathcal{P}(O_i \times R_i) \rightarrow \mathcal{P}(O_{i+1} \times R_{i+1}),$$

and monotonicity becomes $C_{i+1} \subseteq \mathcal{T}_{i \rightarrow i+1}(C_i)$: the successor authority is no greater than what the predecessor permits under the local translation. We treat $\mathcal{T}$ as fixed policy input; synthesizing or verifying such translations is out of scope. Heterogeneity here is semantic rather than merely administrative: different hops may describe authority through different operation/resource vocabularies even within a single trust perimeter.

Lemma 1 (Continuity implies origin binding). *If $PoC(\pi)$ holds, then s_n is causally linked to s_0 through a verifiable relational chain and $C_n \subseteq C_0$.*

Proof. Each $PoR(s_i, s_{i+1})$ links consecutive steps; their composition links s_n to s_0 . Non-expansiveness at each hop gives $C_n \subseteq \dots \subseteq C_0$.

We state the scope explicitly: this paper models the continuity invariant; the concrete (e.g. cryptographic) construction of *PoR* is out of scope and is realized by a companion enforcement architecture. As capability and token models assume their credentials unforgeable, we assume the lineage is unforgeable.

5.1 Authorization Rule

The following rule is evaluated only for a valid continuity chain, i.e. a chain in which every adjacent transition satisfies both $PoR(s_i, s_{i+1})$ and $C_{i+1} \subseteq C_i$.

For the common privilege universe $O \times R$, (o, r) is authorized at the final hop iff:

$$(o, r) \in \bigcap_{i=0}^n C_i.$$

Since the chain is monotonic decreasing:

$$C_n \subseteq C_{n-1} \subseteq \cdots \subseteq C_0,$$

we have:

$$\bigcap_{i=0}^n C_i = C_n.$$

Thus no hop may acquire authority not already present at the origin, and any authority dropped at some hop j is lost irreversibly: by monotonicity it cannot reappear at any later hop.

5.2 The Ontological Shift: Possession vs Continuity

Possession- and capability-based models [2, 3] answer a *spatial* question: does an actor hold authority over an object? Authority is a point in $O \times R$. This model is atemporal: it has no axis on which to record the execution that led to an action.

Proof-of-Relationship supplies the missing axis: it binds an execution step to its causal predecessor. Proof-of-Continuity is built on it — continuity is the transitive composition of relationships — and combines two conditions: monotone restriction of permissions ($C_{i+1} \subseteq C_i$) and causal continuity, via *PoR*, of the execution that carries them.

Authority therefore no longer lives in $O \times R$ alone but in $O \times R \times \mathcal{L}$, where $\mathcal{L}$ is the causal lineage of the execution. Concretely, $\mathcal{L}$ is the set of finite linear causal lineages: a lineage is a sequence $\ell = \langle s_0 \prec s_1 \prec \cdots \prec s_n \rangle$, where $\prec$ denotes causal precedence, so each event is an occurrence of a privilege within a particular lineage. Two actions identical in possession — same principal, same token, same privilege — may be distinct in continuity, because they descend from different causes. A request to *pay 100*, issued once by the user and once by a replay, is indistinguishable under possession, yet the two are distinct executions: one continues an authorizing lineage, the other continues none. Possession is the projection onto $O \times R$ that forgets the lineage axis. Possession is the spatial component of authority; lineage is its temporal component, and a possession-only model is therefore a projection of authority that forgets time. This does not make possession-based models wrong; it makes them incomplete for the multi-hop case, where authority continuation becomes a temporal property. Thus the relevant event is not merely (p, o, r) , or identity plus privilege, but an occurrence of a privilege within a lineage. When the same principal and the same privilege occur in

two different lineages, the identity coordinate does not separate them; the lineage coordinate does.

In the projection language this is the *temporal reading* of the confused deputy of Definition 2: an action spatially valid but temporally invalid — the deputy holds (o, r) , yet $(o, r) \notin C_0$ for the request that caused it — a collision under the projection that forgets $\mathcal{L}$, two distinct executions that possession cannot tell apart. Section 6 gives the operational characterization (the *open passthrough*), which refines Definition 2 with designation.

Relation to capabilities. Capabilities solve the spatial question, and they do more: by fusing designation and authority in the act of invocation, a capability binds authority to a single causal step. In our terms this is *single-hop continuity* — a Proof-of-Relationship. The classical confused deputy is thereby ruled out by capabilities, not merely reduced. What *capability-as-possession* lacks — a held capability applied without occurrence binding — is the extension of that binding across a chain: it may still be exercised outside the lineage that requested it, and in that use falls into lineage-invariant possession. Proof-of-Continuity is exactly this extension: a capability invocation can be viewed as a one-hop instance of the continuity discipline, and PIC generalizes that discipline across multi-hop execution. PIC is therefore a temporal refinement of capability-style authority propagation, not a replacement for it; the temporal dimension is not an optional stronger check but one that possession alone cannot encode.

5.3 Projection: Lineage-Invariant Possession Cannot Individuate Occurrences

We make the projection precise. An *event* is an occurrence exercising a privilege within a lineage,

$$e = (o, r, \ell) \in E := O \times R \times \mathcal{L},$$

and $p : E \rightarrow O \times R$, $p(o, r, \ell) = (o, r)$, is the projection that forgets the lineage.

Definition 6 (Possession policy). *An authorization policy $A : E \rightarrow \{0, 1\}$ is possession-based iff it is invariant under lineage: it factors through p , i.e. there exists $\bar{A} : O \times R \rightarrow \{0, 1\}$ with $A = \bar{A} \circ p$. Equivalently, $A(o, r, \ell) = A(o, r, \ell')$ for all $\ell, \ell' \in \mathcal{L}$.*

Theorem 3 (Lineage-invariant policies cannot individuate occurrences). *Let $e = (o, r, \ell)$ and $e' = (o, r, \ell')$ be two events with the same privilege but*

distinct lineage, $p(e) = p(e')$ and $\ell \neq \ell'$. Then every possession-based policy satisfies $A(e) = A(e')$.

Proof. $A = \bar{A} \circ p$ and $p(e) = p(e')$, hence $A(e) = \bar{A}(p(e)) = \bar{A}(p(e')) = A(e')$.

The confused deputy and the replay are two instances of the same forbidden pair: events sharing (o, r) but differing in lineage, one of which must be authorized and the other denied. In the confused deputy one lineage originates from a context holding (o, r) and the other does not; in a replay one lineage continues an authorizing request and the other continues none. By the theorem, *no* lineage-invariant possession-based policy can separate them.

Theorem 4 (Any resolution reintroduces continuity). *A policy A authorizes e and denies e' , with $p(e) = p(e')$ and $\ell \neq \ell'$, only if A is not lineage-invariant; that is, A must read some function $g : \mathcal{L} \rightarrow X$ of the lineage that is non-constant on the fiber $p^{-1}(o, r)$.*

Proof. If A were lineage-invariant it would factor through p and give $A(e) = A(e')$ by the previous theorem, contradicting $A(e) \neq A(e')$. Hence A depends on the lineage coordinate: some $g(\ell)$ distinguishes e from e' .

Such a g is exactly a continuity mechanism — a proof-of-relationship extract of the lineage. The mitigations deployed within possession-based systems (nonces, idempotency keys, timestamps, audience and context binding) are concrete instances of g : they resolve the confused deputy or the replay only by reintroducing lineage information — that is, continuity — into the decision. Whether the lineage is carried alongside the request or embedded inside the artifact is immaterial: a token that encodes a nonce or a causal context makes A depend on ℓ and is therefore continuity-aware by definition, not possession-based. The distinction is not where the information lives but whether the decision reads it. The limitation is not an engineering failure of possession-based systems; it is a modeling boundary of lineage-invariant authorization: by the first theorem it is structurally unable to individuate occurrences, and every working fix is lineage-sensitive — continuity-aware — rather than lineage-invariant.

Why the possession model is representative. In the abstraction considered here, a bearer token used without explicit context or lineage binding is lineage-invariant: the decision reads the artifact and the requested

privilege, not the causal lineage of the use. A JWT, for instance, asserts that “the holder may read,” not the causal chain that led to it. Such systems often work not because they encode continuity explicitly, but because a perimeter supplies part of it implicitly — mutual TLS, network boundaries, and transport-inferred delegation often stand in for the missing lineage. Once the perimeter dissolves — zero-trust deployments, multi-organisation chains on the public internet, and autonomous AI agents calling tools and services — lineage-invariance becomes visible as a vulnerability at the second hop. The definition of a possession policy above therefore describes deployed systems, not a caricature of them. A trust perimeter may hide the temporal gap by making authority and causality appear co-located, but the model does not depend on crossing such a perimeter: the mismatch can arise whenever authority is selected independently of the execution that caused its use.

5.4 A Trade-off Theorem for Authority Propagation

The projection of Section 5.3 is the technical basis for a trade-off among three properties an authority-propagation system may be asked to satisfy at once.

Definition 7 (Lineage-invariant authorization). *Authorization of an event $e = (o, r, \ell)$ depends only on the possessed authority for (o, r) , not on the lineage ℓ that caused the event.*

Definition 8 (Authority-mixing-capable delegation). *An executor may hold both a request-derived authority context C_0 and an independent authority source A_D , with some (o, r) such that $(o, r) \notin C_0$ and $(o, r) \in A_D$.*

Definition 9 (Confused-deputy safety). *No request whose origin context C_0 lacks (o, r) can cause a later executor to exercise (o, r) using authority not derived from that request.*

Theorem 5 (Possession–delegation–safety trade-off). *No authority-propagation system can simultaneously satisfy lineage-invariant authorization, authority-mixing-capable delegation, and confused-deputy safety.*

Proof. Suppose all three hold. By authority-mixing-capable delegation there is an executor D with a request-derived context C_0 and an independent source A_D , and an (o, r) with $(o, r) \notin C_0$ and $(o, r) \in A_D$. Let a request cause D to process a target requiring (o, r) . Lineage-invariant authorization reads that D possesses A_D granting (o, r) ; but by the projection theorem of Section 5.3, events sharing (o, r) yet differing in lineage receive the same judgment, so

the policy cannot condition on whether A_D belongs to the lineage of this request. The system therefore admits an execution in which the request causes D to exercise (o, r) while $(o, r) \notin C_0$, violating confused-deputy safety — a contradiction.

The theorem is not an attack on possession: it shows that lineage-invariant possession is insufficient for multi-hop authority propagation precisely when authority mixing and confused-deputy safety are both required. The trade-off is resolved only by relinquishing one property — forbidding authority-mixing-capable delegation, accepting confused-deputy executions, or making authorization lineage-sensitive. Proof-of-Continuity takes the third option: it does not remove possession but restores the missing temporal coordinate, making authorization lineage-sensitive so that each authority use is a valid continuation of the context that caused it. The contribution is not to reject possession but to show that possession is the spatial projection of a richer authority relation; the missing coordinate is temporal lineage.

Corollary 1 (Continuity restores confused-deputy safety). *Relinquishing lineage-invariance regains safety: under non-expansive propagation $C_{i+1} \subseteq C_i$, so $C_k \subseteq C_0$, and an $(o, r) \notin C_0$ cannot appear at any later hop. This is the impossibility established in Section 7.*

6 Safety Property

Definition 10 (Origin-bounded authority). *A model enforces origin-bounded authority if every executable privilege at hop n was originally granted at hop 0.*

We refine Definition 2 to separate legitimate delegation from the confused deputy. The distinguishing variable is not the authority source but whether the action’s target is designated by the request and whether the requester was entitled to bring it about.

Definition 11 (Authority to cause). *The origin context C_0 records the authority the requester is entitled to cause, not only to perform directly. A later action belongs to the requester’s lineage iff it realizes an authority in C_0 ; otherwise it is the deputy’s own action, rooted in a distinct origin (the deputy itself) and unconstrained by C_0 .*

Definition 12 (Open passthrough). *A deputy is an open passthrough when, as a consequence of a request, it applies authority to a target designated by*

that request for an action the requester was not entitled to cause: $(o, r) \notin C_0$ yet (o, r) executes because of the request. This is exactly the confused deputy.

These separate three cases. (i) A deputy's own housekeeping — the compiler writing its statistics file — is rooted in the deputy as origin and is unconstrained by C_0 . (ii) Legitimate delegation — a payment gateway moving funds because the user was entitled to cause the payment — realizes an authority in C_0 and is a valid continuation, even though the user does not directly hold the banking authority. (iii) The confused deputy — writing a user-designated file the user could not cause — is an open passthrough, the one case the model forbids. The relevant question is not whether the origin could perform the downstream operation directly, but whether the origin was authorized to *cause* it within this lineage. When successive services use different operation spaces, this is mediated by the policy translation $\mathcal{T}$ introduced above, whose synthesis and verification are left to future work.

Example. If $C_0 = \{(\text{convert}, r)\}$, then (write, r) can never be authorized within the same execution chain: $(\text{write}, r) \notin C_0$ and $C_{i+1} \subseteq C_i$ at every transition, so $(\text{write}, r) \notin C_n$.

7 Impossibility of the Confused Deputy

The confused deputy condition requires an authority mismatch: a privilege absent from the authority context of the request is nevertheless exercised as a consequence of that request.

This is the paper's second impossibility result. Section 5.3 shows that possession *cannot* individuate occurrences; here we show the dual: under continuity the confused deputy is *not a state the model can represent*. The condition is unsatisfiable within any valid PIC execution — an impossibility by construction, at the level of valid model behavior.

Theorem 6 (Confused deputy is impossible under PIC). *In any system satisfying PIC monotonicity $C_{i+1} \subseteq C_i$ and origin-bounded authority $C_n \subseteq C_0$, the confused deputy conditions cannot be jointly satisfied as valid behavior within the model.*

Proof. Immediate from monotonicity: $C_k \subseteq C_0$ for all k , so $(o, r) \notin C_0$ implies $(o, r) \notin C_k$. The privilege mismatch a confused deputy requires is therefore unsatisfiable in a valid PIC execution.

Interpretation. Authority in PIC is not something that a subject merely *has*; it is a continuity property of the execution chain. Since no hop may introduce or reinterpret authority absent at the origin, a confused-deputy execution may still be attempted physically, but it cannot be accepted as valid behavior within the PIC model.

8 Discussion

PoP systems intentionally derive authorization from artifact control, which is often useful and necessary for local authorization decisions. In multi-hop execution, however, artifact control alone does not determine whether a later use is a valid continuation of the authority lineage that caused it. PIC/PoC systems instead propagate only non-expansive subsets of authority along a causally verified execution chain.

Proof of Possession proves artifact control. Proof of Relationship proves execution relationship. Proof of Continuity proves authority continuation. Proof-of-Possession establishes control over an authority artifact at a point in execution; it does not by itself establish that a later use is a valid continuation of the authority lineage that caused it, which is the temporal condition Proof-of-Continuity adds.

Design space. The temporal binding may be realized in two ways: embedded *inside* the authority itself, as attenuated and context-bound capabilities, or maintained *alongside* it as an explicit continuity proof. PIC takes the latter position, treating continuity as a first-class property of the execution chain; the two approaches are complementary rather than exclusive.

The resulting security invariant is simple: no downstream service can perform, as part of a valid execution chain, an operation that was not authorized at the origin. This is a modeling refinement rather than a criticism of existing mechanisms: possession, capabilities, bearer tokens, and proof-of-possession bindings remain useful local authorization tools, while PoC makes explicit the additional temporal condition needed when authority is propagated across an execution lineage.

Authority as a distributed transaction. A multi-hop authority chain is a distributed transaction over authority rather than over data. A possessed artifact enables a local authorization decision at the point of use, which is often necessary for deployability, performance, and loose coupling; but revocation, delegation state, scope changes, and lineage are global authority facts.

Local authorization and global authority consistency can diverge when these facts are not represented at the point of use. PoC addresses this divergence by carrying lineage explicitly and requiring non-expansive authority propagation across the chain. Continuity is therefore not identity propagation. Identity remains useful for authentication, audit, attribution, accountability, and origin creation, but it is insufficient as the propagated authorization primitive. Identifiers are the technical handles by which identities, workloads, roles, or executors are named and bound to evidence; they are necessary for verification and accountability, but an identifier repeated at later hops does not by itself prove that the authority being exercised belongs to the same execution lineage. Identity providers, wallets, OIDC/OIDC4VC-style flows [13, 14], WIMSE/SPIFFE-style workload identity [11, 12], and machine identities remain important for authentication, credential presentation, technical binding, and origin bootstrap. PoC concerns what happens after such a permissioned entity has expressed an intent within its privileges: the authorization question becomes whether later authority uses are valid continuations of the origin authority context C_0 , not whether the same identity or identifier is repeated at each hop. A lineage begins when a permissioned entity expresses an intent within its privileges, thereby creating an origin authority context C_0 . Subsequent hops need not carry, reveal, or reinterpret the origin identity; they must only preserve a non-expansive authority context that remains causally linked to that origin.

Application to AI agents. For the purposes of authority propagation, AI agents can be modeled as autonomous distributed systems; we name them explicitly because the temporal gap is most acute there, but the model treats them as a special case of distributed execution. An agent executes a chain of steps — planning, tool invocations, and calls to downstream services — while holding multiple authority sources at once: its own service credentials, delegated user tokens, and per-tool scopes. The classical local-filename designation reappears as an agent-influenced parameter: a tenant id, document id, role ARN, queue message, cache key, or downstream request target. When the caller can influence this target and the agent satisfies it using authority from a different context, the result is the distributed analogue of the confused deputy. PIC binds each tool call and downstream request to the causal lineage of the request that triggered it, so an agent cannot apply authority that does not belong to the execution that caused the action. For example, an agent asked to summarize a document that contains an indirect prompt injection — one instructing it to delete files, exfiltrate secrets, or

invoke an administrative API — may attempt those actions; under PoC they are valid only if the origin authority context contains the corresponding privileges. If the origin authorizes only $(\text{summarize}, d)$, then (delete, r) or $(\text{exfiltrate}, r)$ cannot be valid continuations of that request.

Related Work

The confused deputy problem was introduced by Hardy [1]. Capability-based access control [2] and the analysis of the confused deputy under capabilities [3] address the spatial dimension of authority — whether an actor holds authority over an object. This paper isolates the orthogonal temporal dimension: whether the authority used belongs to the execution that caused the action. It focuses narrowly on the confused deputy as a mismatch between authority and causality, and compares possession-based authority propagation with provenance-continuity-based authority propagation.

The closest prior work is history-based access control [4], where the authority available to a computation depends on the history of what has executed. We share the premise that authority is a function of execution rather than of static possession. Unlike a retrospective check over past callers (as in stack inspection), PoC is a forward propagation discipline: each hop must be causally related to its predecessor and must not expand the authority it received, including across asynchronous and cross-boundary steps. The contribution here is orthogonal and more specific: we frame the dependence as a *dimension* $\mathcal{L}$ orthogonal to the privilege (o, r) , prove that possession (any policy invariant under $\mathcal{L}$) cannot individuate occurrences, and show that the confused deputy and the replay are the same collision under the projection that forgets $\mathcal{L}$ — so that possession-based mitigations that distinguish such cases must become lineage-sensitive, hence continuity-aware. The object-capability literature [5] analyses the confused deputy as a separation of designation from authority, which capabilities reunify; this is the spatial resolution. Our result adds the temporal axis: in the abstraction of this paper, a capability treated only as a portable held artifact, without occurrence or lineage binding, falls into the lineage-invariant projection and cannot distinguish two occurrences that differ only in lineage.

Bearer tokens, macaroons, and proof-of-possession bindings. OAuth 2.0 includes bearer-token usage as a common authorization pattern [6, 7]. In the abstraction considered here, a bearer token used without explicit context or lineage binding supports an authorization decision that

reads the token and (o, r) , but not the causal lineage of the use: it is lineage-invariant, and thus subject to the first theorem. Macaroons [8] add *caveats* that only attenuate and chain them cryptographically — precisely our monotone attenuation $C_{i+1} \preceq C_i$ along a chain — yet a macaroon remains a bearer credential: absent a caveat that binds the request context, it can be presented outside the lineage that produced it. DPoP [9] and Token Binding [10] bind a token to a key or transport channel, which is exactly a function g of the occurrence in the sense of Theorem 4. The unifying observation is that these mechanisms defeat replay or cross-context misuse precisely to the extent that a caveat or binding makes the decision read ℓ — that is, they reintroduce continuity. They are therefore continuity-aware mechanisms in the sense of this model — reading the lineage coordinate at a single hop or through context binding — rather than purely lineage-invariant possession mechanisms; PIC makes the same lineage dimension global and explicit. Such mechanisms strengthen possession but do not by themselves establish global continuity.

The PIC model is developed independently in [17], where Provenance Identity Continuity is introduced as a general model for distributed execution systems. The present paper extends an earlier version [18], refining it with the relationship/continuity distinction and the spatial/temporal analysis of authority.

9 Limitations and Future Work

This paper presents a minimal linear model of Proof-of-Continuity; several issues are intentionally left outside its scope. First, revocation is not treated as a retroactive operation on an already issued lineage: a practical system must decide whether revocation blocks only future transitions, invalidates outstanding continuations, or triggers lineage-wide cancellation. Second, the model is linear; forks, joins, fan-out, fan-in, and DAG-shaped executions require a partial-order generalization of $\mathcal{L}$. Third, the construction of *PoR* is abstract: concrete implementations may use signatures, hash chains, verifiable credentials [16], secure channels, trusted execution environments, or tamper-evident logs, with implementation-dependent security and performance. Fourth, lineage validation is an implementation concern rather than a limitation of the model: practical realizations need not revalidate an entire history at every hop, and may use checkpoints, summarized proofs, compact encodings, or accumulator-style constructions to keep validation cost bounded. Finally, heterogeneous systems require policy translations $\mathcal{T}$

between local vocabularies, whose synthesis, verification, and governance we leave to future work.

10 Conclusion

This paper introduced Proof-of-Continuity as a minimal temporal model for authority propagation. The core distinction is between possession, which proves control over an authority artifact; relationship, which proves causal linkage between adjacent execution steps; and continuity, which proves that authority has propagated without expansion across the whole lineage. The main result is that lineage-invariant possession cannot distinguish occurrences that share the same privilege but arise from different causes; conversely, under Proof-of-Continuity a privilege absent from the origin authority context cannot appear later in the same valid execution chain. The confused deputy condition is therefore not accepted as valid model behavior — it is precisely the authority/causality mismatch that continuity rules out. More broadly, PoC reframes authority propagation as a temporal property of an execution lineage rather than merely a static property of artifact possession.

Note on the name PIC. In Provenance Identity Continuity, the term Identity refers to the security identity or execution identity of the permissioned entity from which authority is created or anchored at the origin. It does not mean that identity is propagated as the authorization primitive at every hop. The model instead shifts the authorization burden from repeatedly interpreting identity to verifying execution continuity: a causal lineage carrying a non-expansive authority context. Provenance denotes the causal origin and lineage of the execution, Identity denotes the authenticated permissioned origin, and Continuity denotes the preservation of authority across later hops without expansion. Identity and identifiers therefore remain essential for authentication, credential presentation, audit, accountability, and origin bootstrap; PoC concerns the authorization continuity that must hold after that origin has been established.

Acknowledgments

The model, definitions, arguments, and mathematical formalization presented in this paper are the author’s own work. Automated language-assistance tools were used only to accelerate drafting and LaTeX editing, and to help

present the author’s reasoning through clearer exposition, English phrasing, mathematical notation, and formatting.

References

- [1] N. Hardy. The Confused Deputy: Or Why Capabilities Might Have Been Invented. *ACM SIGOPS Operating Systems Review*, 22(4):36–38, 1988.
- [2] J. B. Dennis and E. C. Van Horn. Programming Semantics for Multiprogrammed Computations. *Communications of the ACM*, 9(3):143–155, 1966.
- [3] M. S. Miller, K.-P. Yee, and J. S. Shapiro. Capability Myths Demolished. Technical Report SRL2003-02, Systems Research Laboratory, Johns Hopkins University, 2003.
- [4] M. Abadi and C. Fournet. Access Control Based on Execution History. In *Proceedings of the Network and Distributed System Security Symposium (NDSS)*, 2003.
- [5] M. S. Miller. Robust Composition: Towards a Unified Approach to Access Control and Concurrency Control. PhD thesis, Johns Hopkins University, Baltimore, MD, USA, 2006.
- [6] D. Hardt, editor. The OAuth 2.0 Authorization Framework. RFC 6749, Internet Engineering Task Force, October 2012.
- [7] M. Jones and D. Hardt. The OAuth 2.0 Authorization Framework: Bearer Token Usage. RFC 6750, Internet Engineering Task Force, October 2012.
- [8] A. Birgisson, J. G. Politz, Ú. Erlingsson, A. Taly, M. Vrabie, and M. Lentczner. Macaroons: Cookies with Contextual Caveats for Decentralized Authorization in the Cloud. In *Proceedings of the Network and Distributed System Security Symposium (NDSS)*, 2014.
- [9] D. Fett, B. Campbell, J. Bradley, T. Lodderstedt, M. B. Jones, and D. Waite. OAuth 2.0 Demonstrating Proof of Possession (DPoP). RFC 9449, Internet Engineering Task Force, September 2023.

- [10] A. Popov, M. Nyström, D. Balfanz, and J. Hodges. The Token Binding Protocol Version 1.0. RFC 8471, Internet Engineering Task Force, October 2018.
- [11] IETF WIMSE Working Group. Workload Identity in Multi-System Environments (WIMSE). Internet Engineering Task Force, 2024. <https://datatracker.ietf.org/wg/wimse/>.
- [12] SPIFFE Project. Secure Production Identity Framework for Everyone (SPIFFE) Specification. Cloud Native Computing Foundation, 2017. <https://github.com/spiffe/spiffe>.
- [13] OpenID Foundation. OpenID Connect Core 1.0. 2014. https://openid.net/specs/openid-connect-core-1_0.html.
- [14] OpenID Foundation. OpenID for Verifiable Credentials (OpenID4VC). Specification, 2024. <https://openid.net/sg/openid4vc/>.
- [15] World Wide Web Consortium (W3C). Decentralized Identifiers (DIDs) v1.0. W3C Recommendation, 2022. <https://www.w3.org/TR/did-core/>.
- [16] World Wide Web Consortium (W3C). Verifiable Credentials Data Model v1.1. W3C Recommendation, 2022. <https://www.w3.org/TR/vc-data-model/>.
- [17] N. Gallo. PIC Model — Provenance Identity Continuity for Distributed Execution Systems. Zenodo, 2025. <https://zenodo.org/records/17777421>.
- [18] N. Gallo. Authority Propagation Models: PoP vs PoC and the Confused Deputy Problem. Zenodo, 2025. <https://doi.org/10.5281/zenodo.17833000>.