

Policy Store

Interoperability for Policy Organization, Persistence,
Exchange, and Evaluation

Nicola Gallo

16 July 2026

Status	Draft
Version	0.1
Scope	Policy Store

Abstract

This document explains why interoperability is necessary for the organization, persistence, exchange, and evaluation of authorization policies. Policies are expressed in versioned languages, evaluated by versioned runtimes, and often need to be co-located with the decision point that evaluates them. Organizing, persisting, and transferring them reliably — including over constrained networks — is therefore an architectural requirement.

Contents

1	Context and Problem Statement	2
2	Policy Language and Runtime Structure	2
2.1	Policies as Objects	2
2.2	Logical Grouping	2
2.3	The Manifest	2
3	Policy Types	3
3.1	The Canonical Execution Model	3
3.2	Execution Policies	4
3.3	Semantic Translation Policies	4
3.4	Application Policies	5
3.5	The Limits of the Subject-Centric Model	5
3.6	The Non-Expansion Invariant	5
3.7	Grouping by Policy Type	6
4	Transferability and Persistence	6
4.1	Transfer Only What Changes	6
4.2	Immutable, Verifiable Objects	7
4.3	Everything Is a Transferable Object	7
4.4	Minimal Persistence	7
4.5	Versioning and Auditing	7
5	The Ledger: A Git-like Structure	7
5.1	Structure	8
5.2	Why This Solves the Problem	8
5.3	A Runtime, Not a Bundle	8

1 Context and Problem Statement

Policies are, first of all, languages. Like any language, a policy language is expressive within the context it was designed for; no single language fits every domain, and this is precisely why different types of policy languages exist.

A policy language is evaluated by a runtime. Runtimes are versioned, and policy languages themselves potentially evolve through versions as well: a policy can only be evaluated by a runtime that supports the version of the language the policy is written in. Language version and runtime version must therefore be matched explicitly.

Furthermore, there are different types of policies — not only application policies, as discussed later in this document. Policies often need to be co-located with the decision maker, the Policy Decision Point (PDP). This requires a way to transfer policies efficiently, suitable for low-bandwidth links as well, and protected so that an interrupted transfer can be resumed safely.

For these reasons, policies must be organized, persisted, and made transferable and understandable across systems: this is the interoperability problem this document addresses.

2 Policy Language and Runtime Structure

2.1 Policies as Objects

A policy must be treated, first of all, as an *object* — not as a file. A file carries no guarantees about its own meaning: its interpretation depends on where it happens to be placed, on naming conventions, and on the tooling that reads it. An object, by contrast, is self-describing. To be exchanged between systems, a policy object must carry at least the following properties:

- a stable, unique *identifier*, so that the policy is addressable independently of any storage layout;
- the *language* it is written in and the *language version* it targets;
- its *type* (for example, a policy proper versus the schema the policies are written against);
- its *content*, in a form the declared runtime can evaluate.

Identity and meaning are intrinsic to the object: they travel with it, they do not derive from its position in a directory tree.

2.2 Logical Grouping

For the same reason, policies are not grouped as filesystem trees but as *logical groups*. A group collects the policies that are meant to be managed and evaluated together. The filesystem layout used while authoring is an editing convenience; it is not the semantic structure of the group, and it must not be what a decision point relies on. Within a group, policies are organized under logical *paths* (partitions), and each partition declares its own characteristics — language, schema, runtime. A group is therefore potentially *multi-language*: policies under different paths may be written in different languages and evaluated by different runtimes, while still being organized, versioned, and transferred as a single unit.

2.3 The Manifest

Every group must therefore include a *manifest*: a declaration that states, for each logical path, what type of policies it contains, in which language and language version they are expressed, and which runtime — at which version — is able to evaluate them. The manifest is what allows a consumer, a PDP in particular, to decide *before* loading the group whether it can evaluate it at all, resolving the language-version/runtime-version matching problem stated in the previous section. A non-normative pseudo-JSON manifest:

```

{
  "metadata": {
    "name": "commerce-authz",
    "description": "Authorization model for the commerce domain"
  },
  "runtimes": {
    "cedar-stable": {
      "language": { "name": "cedar", "version": ">=4.0.0" },
      "engine": { "name": "sample-engine", "version": ">=1.2.0" }
    },
    "rego-stable": {
      "language": { "name": "rego", "version": ">=1.0.0" },
      "engine": { "name": "other-engine", "version": ">=0.60.0" }
    }
  },
  "partitions": {
    "/": { "runtime": "cedar-stable", "schema": true },
    "/agents": { "runtime": "cedar-stable", "schema": false },
    "/infra": { "runtime": "rego-stable", "schema": false }
  }
}

```

The structure captures the three concerns separately. The **runtimes** section declares one or more evaluation environments: each entry binds a language and a version range to an engine and its version range, so compatibility is expressed as explicit constraints rather than assumed. The **partitions** section maps every logical path of the group to one declared runtime and states whether that path carries a schema. Because each partition binds to its own runtime, a single group can mix policy languages without ambiguity: a consumer knows, path by path, what it must be able to evaluate. Every partition must reference a declared runtime, and a root partition must exist: the manifest is validatable on its own, and the group is rejected as a whole when the declaration is inconsistent.

Because the manifest is itself an object of the group, the group is self-describing: it can be validated, versioned, and transferred as a single coherent unit — a property the rest of this document builds upon.

3 Policy Types

Authorization is usually discussed as if application policies were the whole story. Looking at an execution end to end, however, three distinct types of policies emerge. To see why, we first need a minimal model of what an execution is.

3.1 The Canonical Execution Model

An execution is nothing more than a *transaction initiated by a root*: a principal expresses an intent, and that intent is carried out by a chain of executors. At any point of the chain there is a previous executor ($n - 1$), a current executor (n), and a next executor ($n + 1$): the execution passes from $n - 1$ through n to reach $n + 1$, always as a consequence of the request that started at the root (Figure 1). The chain is not instantaneous, either: a message sent at time x may be handled by a workload provisioned only at time $x + y$ — the executors need not all exist at the same time.

Each policy type governs a different point of this picture.




Figure 1: The canonical execution model: a request initiated at the root flows through a chain of executors. The workload serving executor $n + 1$ may be provisioned after the message was sent.

3.2 Execution Policies

Execution policies govern the execution chain itself: whether and how step n may reach step $n + 1$, across which paths, under which conditions. They do not answer application questions; they shape the way the transaction is allowed to move.

Consider a security team analyzing a network. Starting from a single request initiated at the root — carrying a single token — the team wants to segment the network (micro-segmentation) and insert policies along the execution paths. Those policies talk about segments, flows, and protocols, not about business entities:

```
# Execution policy (pseudo-code): micro-segmentation
permit execution-flow
  from segment "scanners"
  to   segment "db-tier"
  when origin.request == "sec-team/network-scan"
    and origin.token.scope contains "network:scan"
    and protocol == "tcp/5432"

deny execution-flow otherwise
```

3.3 Semantic Translation Policies

Executor n may live in domain A while executor $n + 1$ lives in domain B , and the two domains generally do not share the same vocabulary. When the execution crosses the boundary, the authority expressed in the terms of A must be translated into the terms of B : this is the job of *semantic translation policies*.

For example, a domain that models contracts may express an authority as **ACCEPT-CONTRACT**; the receiving domain has no such concept, and the closest equivalent requires two of its own privileges:

```
# Translation policy (pseudo-code): domain A -> domain B
translate authority "contract:accept"      # vocabulary of A
  into "user:accept" and "admin:accept"    # vocabulary of B
```

A translation must be meaning-preserving and, crucially, non-expansive: what the translated authority permits in B must not exceed what the original authority permitted in A .

3.4 Application Policies

When the execution finally reaches the application, *application policies* evaluate the request in business terms: may this subject perform this action on this resource, given the schema and the context. This is the policy type most authorization systems focus on — necessary, but only one of the three.

3.5 The Limits of the Subject-Centric Model

Most authorization systems and APIs in use today are built around a single request shape: *may subject S perform action A on resource R in context C ?* It is a reasonable shape, and it has served well: it captures application policies faithfully, and it is where the industry has rightly concentrated its effort.

Once executions are looked at end to end, however, questions emerge that this shape was never designed to answer. Execution policies do not speak about subjects and resources: they speak about steps, paths, and flows. Translation policies do not answer an access question at all: they map authority between vocabularies. If a translation policy must be evaluated when the execution crosses a domain boundary, what evaluates it — a second, ad hoc decision point? If an execution policy must be validated inside the network path, does the network now host an application-style PDP?

These are not shortcomings of any particular design: they are problems we all share, surfaced by distributed execution and autonomous agents. The observation they lead to is simple: both the store and the evaluation model benefit from being aware of policy types, rather than assuming that the application type is the whole picture.

3.6 The Non-Expansion Invariant

The root has an *origin*: the authority selected when the intent was expressed. Whatever the policies decide — execution, translation, or application — the execution must never end up holding more authority than that origin. This invariant applies to *all* policy types.

The risk is concrete because policies are chained along the execution: governance effects and application effects concatenate step after step. Suppose one policy in the chain permits more than necessary. If its result is simply accepted, the execution acquires authority it never received from its origin. The rule must therefore be that at every step the effective authority is the *intersection* of what the policies grant with the authority the step received at entry:

$$A_{\text{effective}} = A_{\text{granted}} \cap A_{\text{entry}}$$

The symbol $\cap$ simply reads “what appears in both”: a privilege is effective only if the policies grant it *and* the authority received at entry already contained it. Anything a policy grants beyond the entry authority is clipped; without this rule, authority expands.

This is zero trust taken to its full consequence. Zero trust is commonly read as *never trust, always verify*, applied to identities and connections; applied to authority, it means that every boundary is a point of *attenuation*, never of amplification. Segmentation at the execution level narrows where the transaction may flow; translation at the domain level maps authority without enlarging it; evaluation at the application level grants at most what the origin already allowed. Attenuation applies at every level — execution, domain, application — and the intersection rule is what makes it hold end to end.

Figure 2 shows how an expansion can look perfectly legitimate. Two independent executions (*lineages*) run through the same infrastructure: each one, taken alone, only ever attenuates its authority. A defect then composes authority *across* the lineages, and the final executor of lineage 1 ends up holding *read all* — a privilege its own origin never granted, since it started from *read foo*. Every step looked locally valid; the resulting security state is not.

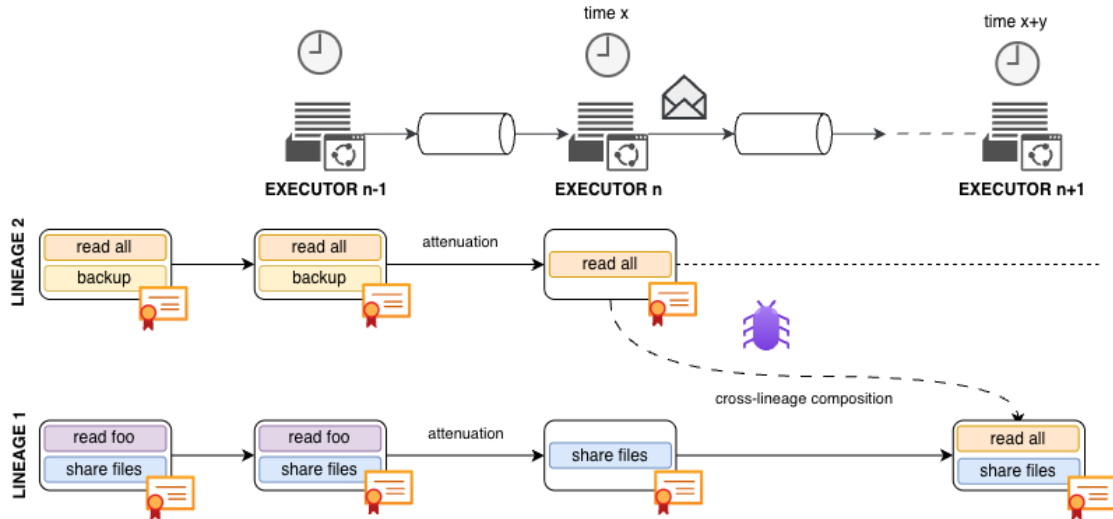


Figure 2: Authority expansion through cross-lineage composition: each lineage only attenuates its own authority, yet a defect merges authority across lineages and the final step holds more than its origin ever granted.

3.7 Grouping by Policy Type

The three types answer different questions, are evaluated at different points of the chain, and are best expressed in different languages: a language designed for network flows is a poor fit for business rules, and vice versa. This is exactly why the manifest of the previous section must allow a group to be organized *per policy type*, with each partition free to declare its own language and runtime:

```
"partitions": {
  "/execution": { "runtime": "netlang-stable", "schema": false },
  "/translation": { "runtime": "maplang-stable", "schema": false },
  "/app": { "runtime": "cedar-stable", "schema": true }
}
```

Note. `netlang-stable` and `maplang-stable` are deliberately fictional names: this is an example only, and no specific language is implied. Execution and translation policies may be expressed in an existing policy language (such as Rego or a JSON-based DSL) or in a purpose-built custom language — the model does not prescribe which language fills which role; it only requires that each partition declares its own.

One group, one coherent unit of governance — several policy types, each in the language that fits it best.

4 Transferability and Persistence

Exchanging policies as described so far imposes a set of properties on how they travel and how they are stored. These properties are independent of any specific storage or network technology.

4.1 Transfer Only What Changes

Policy distribution must work over low-bandwidth links, so exchange must happen only when necessary and must carry only what is necessary. If one policy changed out of ten thousand, only that one travels — never the whole group. This requires being able to tell *what* changed

without comparing contents byte by byte, which is a consequence of how objects are identified, as discussed next.

4.2 Immutable, Verifiable Objects

Every object is *immutable*: a change never modifies an object, it produces a new one. And every object is *verifiable*: the receiver must be able to check, from the object alone, that its content was not altered in transit or at rest.

Both properties follow from *content addressing*¹: the identifier of an object is derived from its content itself. Such an identifier is unique *across* systems — two independent parties compute the same identifier for the same content, with no registry or coordination — and verification reduces to recomputing it. Identical content is also recognized as identical everywhere, which is exactly what makes minimal transfer possible: two sides can agree on what the other is missing by exchanging identifiers, not contents.

4.3 Everything Is a Transferable Object

With identity settled, the picture becomes uniform: policies, schemas, manifests, and the groups that contain them are all objects. A group is transferred as a set of objects and *recomposed* at the destination by following identifiers. If a transfer is interrupted, it can resume safely: objects already received are immutable and verifiable, so storing them again is harmless and nothing needs to be re-negotiated from scratch.

4.4 Minimal Persistence

Persistence must require nothing more than an associative map from identifier to bytes. A directory on disk, an object store such as S3, or an in-memory dictionary are all sufficient backends: no database schema, no query engine, no storage-specific format. This keeps the model deployable everywhere the policies need to go — from cloud services to edge nodes.

4.5 Versioning and Auditing

Finally, the history of a group must be first-class: every change produces a new version, previous versions remain reconstructible, and it must always be possible to answer *which policies were in force at a given point* — a governance and auditing requirement, not a convenience.

5 The Ledger: A Git-like Structure

The properties of the previous section — content-addressed immutable objects, delta transfer, key-value persistence, versioned history — do not need to be invented. A battle-tested system built exactly on them has existed for two decades: Git.

Adopting Git as-is, however, would carry more than the problem needs. It brings along working trees, branching and merging workflows, and hooks; as a general-purpose content tracker it enforces no structure — anything can be committed, and nothing guarantees the invariants of a policy group (a manifest, partitions, one policy per object); and it was not designed for the evaluation path, where a decision point must materialize a precise version of a group quickly, at every decision. A more fitting approach is a *git-like* object model, specialized for policies: battle-tested ideas, without reinventing the wheel and without inheriting what does not serve. We call the resulting structure a *ledger*.

¹<https://docs.ipfs.tech/concepts/content-addressing/>

5.1 Structure

A ledger is an append-only, content-addressed history of one policy group, built from three kinds of objects:

- a *blob* is a leaf object: a policy, a schema, or the manifest;
- a *tree* represents one partition of the group: the list of its entries, each with its name, type, and language metadata;
- a *commit* represents one version of the whole group: it references the manifest blob, one tree per partition, and the predecessor commit.

Every object is serialized in a canonical form and identified by the content-derived identifier of its bytes; the ledger itself is just a *ref*: a pointer to the head commit.

```
commit #C2
predecessor -> commit #C1
manifest    -> blob #M
/execution  -> tree #T1 -> blob #P1, blob #P2
/app        -> tree #T2 -> blob #P3, blob #S (schema)
```

5.2 Why This Solves the Problem

Each requirement stated in this document maps directly onto the ledger:

- **Structure by construction.** A commit must reference a valid manifest, and its trees mirror the declared partitions: a malformed group is rejected at the object level, not by convention.
- **Minimal transfer.** Unchanged content keeps the same identifier, so unchanged partitions share their trees and blobs across versions. Synchronizing two ledgers reduces to comparing refs, computing the missing commits, and transferring only the objects that the other side does not already have — the one-changed-file-in-ten-thousand case transfers one blob, one or two trees, and one commit.
- **Resumable, low-bandwidth exchange.** Objects are verified on receipt by recomputing their identifier, and storing an object twice is idempotent: an interrupted synchronization resumes where it stopped.
- **Minimal persistence.** All objects, whatever their kind, live in any identifier-to-bytes store.
- **Versioning and auditing.** The predecessor chain is the complete history: any past version of the group can be reconstructed exactly, and a decision can be tied to the precise commit it was evaluated against.
- **Evaluation-friendly.** A decision point materializes one commit — manifest, trees, blobs — directly from the store, with no working tree and no checkout.

5.3 A Runtime, Not a Bundle

A natural first answer to distribution is a compressed archive: package the files, ship the bundle to the decision point, unpack, load. Archives are simple and familiar, and for occasional deployments they can be enough. They give up, however, most of what the object model provides: filesystem semantics return in place of self-describing objects, the whole store travels instead of a delta, and there is no per-object verification, no resumable exchange, no history at the destination. Policies end up treated as static configuration to be installed rather than as governed, versioned objects.

Policies are languages, as stated at the outset — and at the point of decision they behave like code: they are brought onto the consumer and executed locally, as if they were libraries loaded into the process. What sits at the consumer is therefore not an unpacker of bundles but a

runtime: a component that synchronizes its local ledger with the policy store and evaluates the synchronized version in place, exactly as a language runtime loads and runs a versioned library.

This is most visible at the edge. A device on an intermittent, low-bandwidth link cannot re-download an archive at every change: it must *synchronize*. It aligns its local ledger when connectivity allows — transferring only the missing objects, resuming where it stopped — and keeps deciding locally against the last verified commit even while disconnected.

Interoperability, then, does not stop at the store format: it extends to the runtime. Any conforming runtime must be able to synchronize any conforming ledger, read its manifest, and evaluate the partitions whose declared languages it supports — so that the same policy group yields the same decisions in the cloud, in a sidecar, or on a device.